

SATURN 2010

Sixth Annual SEI Architecture Technology
User Network Conference

ARCHITECTING FOR CHANGE

May 17-21, 2010
Minneapolis, Minnesota



SAVE THE DATE



SHARE YOUR EXPERIENCE



Software Engineering Institute

Carnegie Mellon

Cloud Computing
November 12, 2009 — SEI Webinar
© 2009 Carnegie Mellon University



SEPGSM 2010
NORTH AMERICA

Perform at a Higher Level

SEPG North America 2010

March 22–25, 2010
Savannah, Georgia



www.sei.cmu.edu/sepg/na/2010

SEPG is the premier, global conference series on
software and systems process management



Software Engineering Institute

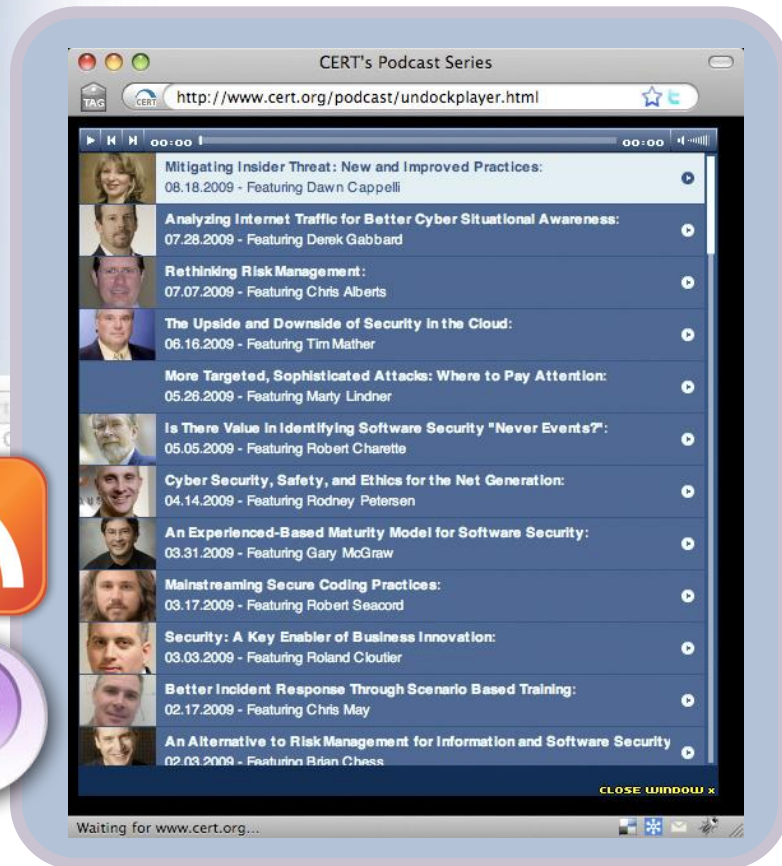
Carnegie Mellon

Cloud Computing
November 12, 2009 — SEI Webinar
© 2009 Carnegie Mellon University



CERT's Podcast Series: Security for Business Leaders

▶ www.cert.org/podcast/





Want a Closer Connection to the SEI?

Become an SEI Member!

▶ www.sei.cmu.edu/membership



Software Engineering Institute

Carnegie Mellon

Cloud Computing
November 12, 2009 — SEI Webinar
© 2009 Carnegie Mellon University



Do you have the knowledge you need?

SEI Training



www.sei.cmu.edu/training



Software Engineering Institute

Carnegie Mellon

Cloud Computing
November 12, 2009 — SEI Webinar
© 2009 Carnegie Mellon University

Cloud Computing

Grace A. Lewis

Research, Technology and Systems Solutions (RTSS) Program
System of Systems Practice (SoSP) Initiative

SEI Webinar
November 12, 2009



Polling Question 1

How did you hear about this webinar?

- Email invitation from the SEI
- SEI Website
- Website with webinar calendar (ie www.webinar-directory.com)
- Social Media site (LinkedIn, Twitter)
- Other



Agenda

Cloud Computing Defined

Types of Cloud

Drivers for Cloud Computing Adoption

Barriers to Cloud Computing

Relationship to Other Technologies

Summary



Polling Question 2

Has your organization adopted cloud computing?

- ☐ No
- ☐ No, but we have plans to do so
- ☐ Just starting
- ☐ Yes



Cloud Computing

“A large-scale distributed computing paradigm that is driven by economies of scale, in which a pool of abstracted, virtualized, dynamically-scalable, managed computing power, storage, platforms, and services are delivered on demand to external customers over the Internet.”¹

Cloud Computing is a distributed computing paradigm that focuses on providing a wide range of users with distributed access to virtualized hardware and/or software infrastructure over the Internet.

¹ I. Foster, Y. Zhau, R. Ioan, and S. Lu. “Cloud Computing and Grid Computing : 360-Degree Compared.” Grid Computing Environments Workshop, 2008.



Cloud Computing — Core Concepts

Distributed computing paradigm

- Large-scale distributed system

Resource-based and service-based

- Provides access to a resource and/or service

Elastic and on demand

- Has the ability to scale up or down based on demand

Economies of scale

- Contains a large number of nodes and therefore reduces the overall cost of resource acquisition and operation

Utility payment models

- Based on pay-per-use model

Based on existing technologies

- Leverages existing technologies such as virtualization, service-orientation and grid computing



Cloud Computing Types

Software-as-a-Service (SaaS)

Platform-as-a-Service (PaaS)

Infrastructure-as-a-Service (IaaS)

Based on Type of Capability

Public Cloud

Private Cloud

Based on Access



Infrastructure-as-a-Service (IaaS)

Generic computational infrastructure available over the Internet, e.g. compute, storage, etc.

Allows organizations and developers to extend their IT infrastructure on an on-demand basis

Examples

- Amazon Elastic Compute Cloud (EC2)
 - Provides users with a special virtual machine (AMI) that can be deployed and run on the EC2 infrastructure
- Amazon Simple Storage Solution (S3)
 - Provides users with access to dynamically scalable storage resources
- IBM Computing on Demand (CoD)
 - Provides users with access to highly configurable servers plus value-added services such as data storage
- Microsoft Live Mesh
 - Provides users with access to a distributed file system; targeted at individual use



Platform-as-a-Service (PaaS)

Application development platforms, e.g. containers to host Java components

Allows developers to leverage the resources of established organizations to create and host applications of a larger scale than an individual or small business would be able to handle

Examples

- Google App Engine
 - Provides users a complete development stack and allows them to run their applications on Google's infrastructure
- Yahoo! Open Strategy (Y!OS)
 - Provides users with a means of developing web applications on top of the existing Yahoo! platform, and in doing so leveraging a significant portion of the Yahoo! resources
- Akamai EdgePlatform
 - Provides a large distributed computing platform on which organizations can deploy their web applications; large focus on analysis and monitoring
- Microsoft Azure Services Platform
 - Provides users with on-demand compute and storage services as well as a development platform based on Windows Azure



Software-as-a-Service (SaaS)

Application-specific capabilities, e.g., service that provides customer management

Allows organizations and developers to use business-specific capabilities developed by third parties

Examples

- Force.com
 - From salesforce.com (SaaS leader), provides enterprise users a platform to build and run applications and components bought from AppExchange or custom applications
- Zoho
 - Provides a large suite of web-based applications, mostly for enterprise use



Cloud Computing Types — Based on Access

Public

- Offered as a service, usually over an Internet connection
- Typically charge a monthly usage fee
- Users can scale on-demand and do not need to purchase hardware
- Service providers manage the infrastructure and pool resources into capacity required by consumers

Private

- Deployed inside the firewall and managed by the user organization
- User organization owns the software and hardware running in the cloud
- User organization manages the cloud and provides cloud resources
- Resources typically not shared outside the organization and full control is retained by the organization

PERSPECTIVE

**Cloud
Consumer**

**Cloud
Provider**



Development Tools for the Cloud Provider

3tera

- Provides developers with tools to build their own cloud computing infrastructures

Eucalyptus Systems

- Provides an open-source application that can be used to implement a cloud computing environment on a datacenter
- Trying to establish an open set of standards for cloud computing



Polling Question 3

In your opinion, what is the main driver for Cloud Computing adoption?

- Reduced costs
- Reduced risk
- Higher scalability
- Other
- Don't Know



Drivers for Cloud Computing Adoption

Scalability	Users have access to a large amount of resources that scale based on user demand
Elasticity	The environment transparently manages a user's resource utilization based on dynamically changing needs
Virtualization	Each user has a single view of the available resources, independently of how they are arranged in terms of physical devices
Cost	The pay-per-usage model allows an organization to only pay for the resources they need with basically no investment in the physical resources available in the cloud. There are no infrastructure maintenance or upgrade costs
Mobility	Users have the ability for the user to access data and applications from around the globe
Collaboration	Users are starting to see the cloud as a way to work simultaneously on common data and information
Risk Reduction	Users can use the cloud to test ideas and concepts before making major investments in technology



Polling Question 4

In your opinion, what is the main barrier to cloud computing adoption?

- Security
- Lack of control
- Reliability
- Performance
- Other



Barriers for Cloud Computing Adoption

Security	The key concern is data privacy: users do not have control of or know where their data is being stored
Interoperability	A universal set of standards and/or interfaces have not yet been defined, resulting in a significant risk of vendor lock-in
Control	The amount of control that the user has over the cloud environment varies greatly
Performance	All access to the cloud is done via the internet, introducing latency into every communication between the user and the environment
Reliability	Many existing cloud infrastructures leverage commodity hardware that is known to fail unexpectedly
Platform or Language Specificity	Some cloud environments provide support for specific platforms and languages only

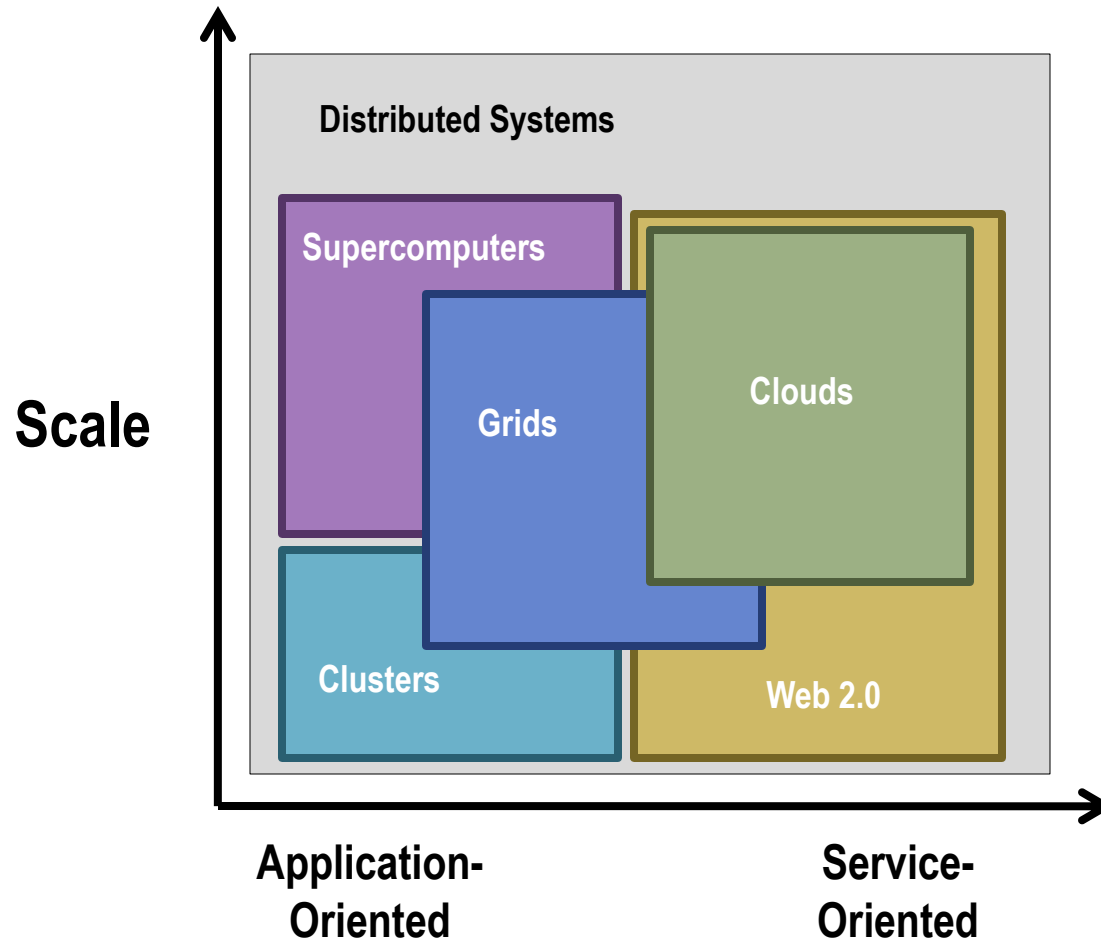


Other Criteria for Cloud Computing Adoption

Feedback	The ability for the user to obtain the statistics on usage levels
Usability	The ease with which the user is able to configure and operate virtual resources



Relationship to Other Technologies



Source: I. Foster, Y. Zhau, R. Ioan, and S. Lu. "Cloud Computing and Grid Computing : 360-Degree Compared." Grid Computing Environments Workshop, 2008.



Relationship to Other Technologies: Service-Oriented Architecture

Service-oriented architecture (SOA)

- A way of designing, developing, deploying, and managing systems characterized by coarse-grained services that represent reusable functionality
- Service consumers compose applications or systems using the functionality provided by these services through standard interfaces.

The services in a cloud can be defined as services in a SOA context.

- As an example, some of the cloud environments presented earlier offer web service interfaces to their services (one specific implementation of SOA)

From an architectural perspective

- A cloud infrastructure could be built on top on an SOA infrastructure by adding a layer or virtualization and self-provisioning
- A service layer could be added on top of cloud resources
- Some even consider SOA adoption a pre-requisite for cloud computing



Relationship to Other Technologies: Software as a Service (SaaS)

Model of software deployment in which a provider licenses an application to customers for use as a service on demand

Multiple Levels

- Level 1: An application is specifically run for one customer at an SaaS provider, similar to the traditional ASP (application server provider) model
- Level 2: The SaaS application is customizable via configuration and one instance of the application serves only one customer.
- Level 3: The SaaS application is customizable and a single instance of the SaaS application serves multiple tenants.
- Level 4: The SaaS application is developed as a single instance multi-tenant application and several instances are run in a load-balanced server farm.

Levels 3 and 4 of SaaS fall under the definition of cloud computing.

³ F. Chong and G. Carraro. Building Distributed Applications Architecture Strategies for Catching the Long Tail. MSDN architecture center, <http://msdn2.microsoft.com/enus/library/aa479069.aspx>, 2006.



Relationship to Other Technologies: Grid Computing

A grid is “a system that uses open, general-purpose protocols to federate distributed resources and to deliver better-than-best-effort qualities of service”¹

Although the distinction with cloud computing is not clear, one differentiator is that grid computing relates exclusively to infrastructure services

- A grid infrastructure provides a set of abstractions and interfaces for access to, and management of, shared resources

Related terms: utility computing, on-demand computing

¹ Foster, I. 2002. What is the grid? A three-point checklist; <http://www-fp.mcs.anl.gov/~foster/Articles/WhatIsTheGrid.pdf>

² Foster, I., Kesselman, C., Nick, J.M., and Tuecke, S. Grid services for distributed systems integration. IEEE Computer 35 (6): 37-46.



Summary

Cloud Computing is in essence an **economic model**

- It is a different way to acquire and manage IT resources

There are multiple cloud providers—**the cloud is real**

- Currently most cloud consumers are small enterprises
- Large enterprises are exploring private clouds
- The number of providers will most probably grow as people start seeing greater savings and work to reduce adoption barriers

Cloud Computing adoption requires **cost/benefit/risk analysis** to determine

- What resources to move to the cloud (if any)
- What situations warrant use of cloud resources, even for one-time situations
- Implementation of private clouds vs. usage of public clouds
- What risks are associated with using resources on the cloud



Contact Information

Grace A. Lewis


Research, Technology and Systems Solutions (RTSS) Program
System of Systems Practice (SoSP) Initiative

Software Engineering Institute
4500 Fifth Avenue
Pittsburgh, PA 15213-2612
USA

Phone: +1 412-268-5851

Email: glewis@sei.cmu.edu

WWW: <http://www.sei.cmu.edu/about/people/glewis.cfm>



Environment and Tool References

- 3tera: <http://www.3tera.com/>
- Akamai EdgePlatform: <http://www.akamai.com/html/technology/edgeplatform.html>
- Amazon Elastic Compute Cloud (EC2): <http://aws.amazon.com/ec2/>
- Amazon Simple Storage Solution (S3): <http://aws.amazon.com/s3/>
- Eucalyptus Systems: <http://www.eucalyptus.com/>
- Force.com: <http://www.salesforce.com/platform/>
- Google App Engine: <http://code.google.com/appengine/>
- IBM Computing On Demand: <http://www-03.ibm.com/systems/deepcomputing/cod/>
- Microsoft Azure Services Platform: <http://www.microsoft.com/azure/>
- Microsoft Live Mesh: <http://www.mesh.com/>
- Yahoo! Open Strategy (Y!OS): <http://developer.yahoo.com/yos/intro/>
- Zoho: <http://www.zoho.com/>



Selected Literature References

- Drdogmus, H. Cloud Computing: Does Nirvana Hide Behind the Nebula? IEEE Software: March/April 2009.
- Foster, I., Y. Zhau, R. Ioan, and S. Lu. Cloud Computing and Grid Computing 360-Degree Compared. Grid Computing Environments Workshop, 2008.2.
- Grossman, R. The Case for Cloud Computing. ITPro: March/April 2009.3.
Hayes, B. Cloud Computing. Communications of the ACM. 51(7): July 2008, p 9-11.4.
- Lin, G., F. David, Z. Jinzy, and D. Glenn. Cloud Computing: IT as a Service. IT Pro: March/April 2009, p 10-13.5.
- McEvoy, G. and B. Schulze. Using Clouds to address Grid Limitations. MGC, December 2008.6.
- Wang, L, J. Tao, M. Kunze, A. Castellanos, D. Kramer, W. Karl. Scientific Cloud Computing: Early Definition and Experience. The 10th IEEE International Conference on High Performance Computing and Communications. 2008.7.
- Weiss, A. Computing in the Clouds. December 2007.



NO WARRANTY

THIS MATERIAL OF CARNEGIE MELLON UNIVERSITY AND ITS SOFTWARE ENGINEERING INSTITUTE IS FURNISHED ON AN "AS-IS" BASIS. CARNEGIE MELLON UNIVERSITY MAKES NO WARRANTIES OF ANY KIND, EITHER EXPRESSED OR IMPLIED, AS TO ANY MATTER INCLUDING, BUT NOT LIMITED TO, WARRANTY OF FITNESS FOR PURPOSE OR MERCHANTABILITY, EXCLUSIVITY, OR RESULTS OBTAINED FROM USE OF THE MATERIAL. CARNEGIE MELLON UNIVERSITY DOES NOT MAKE ANY WARRANTY OF ANY KIND WITH RESPECT TO FREEDOM FROM PATENT, TRADEMARK, OR COPYRIGHT INFRINGEMENT.

Use of any trademarks in this presentation is not intended in any way to infringe on the rights of the trademark holder.

This Presentation may be reproduced in its entirety, without modification, and freely distributed in written or electronic form without requesting formal permission. Permission is required for any other use. Requests for permission should be directed to the Software Engineering Institute at permission@sei.cmu.edu.

This work was created in the performance of Federal Government Contract Number FA8721-05-C-0003 with Carnegie Mellon University for the operation of the Software Engineering Institute, a federally funded research and development center. The Government of the United States has a royalty-free government-purpose license to use, duplicate, or disclose the work, in whole or in part and in any manner, and to have or permit others to do so, for government purposes pursuant to the copyright license under the clause at 252.227-7013.

